

Structs and Arrays

Where We Left Off

A variable is a **symbolic name** for an address

The **contents** of that address is the value

We used atomic types:

i32, u8, f64, bool, char

int32_t, uint8_t, double, bool, char

One name, one value, one location

Real Data Comes in Groups

A single pressure reading is really **three** numbers

- the raw ADC count from the sensor
- the pressure in psi we computed
- the time the sample was taken

They are three variables that belong together

The language should let allow to group them efficiently

Loose Variables Do Not Scale

```
uint32_t adc_counts_1 ;  
float    psi_1 ;  
uint64_t timestamp_1 ;
```

```
uint32_t adc_counts_2 ;  
float    psi_2 ;  
uint64_t timestamp_2 ;
```

Nothing ties them together
The compiler cannot help you keep them
straight

The Struct

A **struct** groups variables into a **new type**
<T>

The grouped variables are called **members**
or **fields**

They are stored in one **consecutive block**
of memory

One address locates the whole group

Same idea in C, C++, and Rust

Also known as a record

Defining a struct in C

```
struct sample {  
    uint32_t  adc_counts ;  
    float     psi ;  
    uint64_t  timestamp_ms ;  
} ;
```

This defines a **type**, not a variable

No memory is allocated yet

Declaring the variable in C

```
struct sample my_sample ;
```

The keyword **struct** is **required** in C

This is an instance of type **struct** sample

<T> is <struct sample>

```
struct sample sample2 ;
```

sample2 is a second instance of the same new type

Typedefs

```
struct sample {  
    uint32_t  adc_counts ;  
    float      psi ;  
    uint64_t  timestamp_ms ;  
} ;
```

```
typedef struct sample sample_t ;  
sample_t my_sample ;  
sample_t sample2 ;
```

Now memory exists: sizeof(struct sample)
bytes

Better typedef

Even better: use **typedef** with the **struct** definition:

```
typedef struct sample {  
    uint32_t  adc_counts ;  
    float      psi ;  
    uint64_t  timestamp_ms ;  
} sample_t ;
```

Variable definitions again are:

```
sample_t my_sample ;
```

```
sample_t sample2 ;
```

Now memory exists: sizeof(sample) bytes

Defining a struct in C++

Same syntax as C:

```
struct sample { ... } ;  
struct sample my_sample ;
```

The struct keyword is **not required** when declaring

C++ also has **class**

struct members are **public** by default

class members are **private** by default

Otherwise functionally equivalent

Defining a struct in Rust

```
struct Sample {  
    adc_counts : u32,  
    psi : f32,  
    timestamp_ms : u64,  
}
```

Fields are separated by **commas**

No semicolon after the closing brace

Type names are **CamelCase**, fields are **snake_case**

Programming Case Styles

Case Style	Format /Rule	Example	
Common Use			
Camel Case and functions (camelCase) JavaScript, Java, C#	First letter lowercase, next words capitalized, no spaces	userId	Variables in
Upper Camel Case Rust (UpperCamelCase)	Every word capitalized no spaces	UserId	Types in
Pascal Case in C#, (PascalCase)	Every word starts with a capital letter, no spaces	UserId	Classes and types Java, TypeScript
Snake Case and functions in Python, Ruby	All lowercase, words (snake_case) separated by underscores (_)	user_id	Variables
Screaming Snake Constants in Python, (SCREAMING_SNAKE) macros C, C++	All uppercase, words separated by underscores (_)	USER_ID	Java,

Defining a struct in Rust

```
struct Sample {  
    adc_counts : u32,  
    psi : f32,  
    timestamp_ms : u64,  
}
```

Fields are separated by **commas**

No semicolon after the closing brace

Type names are **CamelCase**, fields are **snake_case**

You can turn it off for a specific struct by
using **`#[allow(non_camel_case_types)]`**

Not recommended to turn off the standard
Makes it easier to read others code

Rust Programming Guidelines

In general, Rust tends to use UpperCamelCase for "type-level" constructs (types and traits) and snake_case for "value-level" constructs. More precisely:

Item	Convention
Crates	unclear
Modules	snake_case
Types	UpperCamelCase
Traits	UpperCamelCase
Enum variants	UpperCamelCase
Functions	snake_case
Methods	snake_case
General constructors	new or with_more_details
Conversion constructors	from_some_other_type
Macros	snake_case!
Local variables	snake_case
Statics	SCREAMING_SNAKE_CASE
Constants	SCREAMING_SNAKE_CASE
Type parameters	concise UpperCamelCase, usually single uppercase letter: T
Lifetimes	short lowercase, usually a single letter: 'a, 'de, 'src
Features	unclear but see C-FEATURE

Giving it a value: C and C++

```
struct sample my_sample = { 1456, 50.0f, 12345 } ;
```

Order matters! Better, name the members:

```
struct sample my_sample = {  
    .adc_counts    = 1456,  
    .psi           = 50.0f,  
    .timestamp_ms  = 12345  
} ;
```

But we still need this definition for the type to precede it:

```
typedef struct sample {  
    uint32_t    adc_counts ;  
    float       psi ;  
    uint64_t    timestamp_ms ;  
};
```

Giving it a value: Rust

```
let my_sample = Sample {  
    adc_counts    : 1456,  
    psi           : 50.0,  
    timestamp_ms  : 12345,  
} ;
```

Every field must be given a value

The compiler will not let you forget one

Use **let mut** if you intend to change it later

Reaching a Member: the dot

C and C++:

```
my_sample.psi = 50.0f ;  
printf("%f\n", my_sample.psi) ;
```

Rust:

```
my_sample.psi = 50.0 ;  
println!("{}", my_sample.psi) ;
```


The **dot operator** is the same in all three

A struct in Memory

The struct has **one base address** :
&my_sample

Each member sits at a fixed **offset** from that base

address of member = base address +
offset



Size and Offset

The compiler knows every offset at **compile time**

C and C++:

```
sizeof(struct sample)
#include <stddef.h>
offsetof(struct sample, psi)
```

Rust:

```
std::mem::size_of::<Sample>()
```

Print these in your own program - do not guess

Padding and Alignment

Hardware reads memory fastest on **natural boundaries**

a 4 byte value wants an address divisible by 4

an 8 byte value wants an address divisible by 8

The compiler inserts **padding bytes** to make this true

So `sizeof(struct)` can be **larger** than the sum of its parts

Let's look at `mempractice` again

Padding Costs Real Memory

```
struct bad {  
    uint8_t    flag ;  
    uint64_t   stamp ;  
    uint8_t    id ;  
} ;
```



sizeof is 24 bytes: 10 of data, 14 of padding

Order Your Members

Rule of thumb: declare members **largest to smallest**

```
struct good {  
    uint64_t  stamp ;  
    uint8_t   flag  ;  
    uint8_t   id    ;  
} ;
```

sizeof is now **16** bytes, not 24



offset
8

Controlling the Layout

C and C++ (compiler specific):

```
#pragma pack(1)  
__attribute__((packed))
```

Rust:

```
#[repr(C)]  
struct Sample { ... }
```

Rust may **reorder** fields unless you say
#[repr(C)]

Pointers to structs

C and C++:

```
struct sample *p = &my_sample ;  
p->psi = 50.0f ;           same as (*p).psi
```

Rust:

```
let p = &my_sample ;  
p.psi           Rust dereferences for you
```

There is **no -> operator** in Rust

Pointers

Pointers are variables that are interpreted as addresses

It is an indirection to the memory of interest

We use (variable) to denote the contents of an address

A pointer is the address of the contents

pointer to a variable $\leq \& \text{ (variable)}$

If we wish to get the contents pointed by the pointer, we dereference the pointer

contents pointed to $\leq ((\text{pointer}))$

Handles

A handle is a pointer to a pointer

Handle $\leftarrow$ & & (variable)

To get the contents pointed by the handle, we need to dereference twice: handle $\rightarrow$ pointer $\rightarrow$ contents

contents $\leftarrow$ (((handle)))

See mempractice tool in pointer mode

Structs inside structs

```
struct reading {  
    struct sample  sample ;  
    uint8_t        sensor_id ;  
} ;
```

Reach in with more dots:
my_reading.sample.psi

Still one block of memory, still fixed offsets

Build complex objects through hierarchy

Passing a struct to a function

By value copies **every byte**:

```
void show(struct sample s) ;
```

Pass a pointer or reference instead:

```
void show(const struct sample *s) ;
```

```
void show(const sample &s) ;    C++
```

```
fn show(s : &Sample) { ... } Rust
```

In Rust, passing by value **moves** ownership

Many of the Same Thing

One sample is not enough

We want the last **100** samples

Copy and paste 100 variable definitions?

```
my_sample_1, my_sample_2, ...  
my_sample_100
```

There must be a better way

The Array

An **array** is N variables of the **same** type

Stored **consecutively** in memory

One name, one base address

Individual elements selected by an **index**

The index starts at **0**

Declaring arrays in C and C++

```
uint32_t counts[100];  
float     psi[100] = { 0.0f } ;
```

The length is part of the declaration

C++ also offers a safer array:

```
#include <array>  
std::array<uint32_t,100> counts ;
```

Declaring arrays in Rust

```
let counts : [u32 ; 100] = [0 ; 100] ;
```

```
let mut psi : [f32 ; 100] = [0.0 ; 100] ;
```

The type is [type ; length]

The **length is part of the type**

[0 ; 100] means "the value 0, one hundred times"

Indexing starts at zero

counts[0] the **first** element

counts[99] the **last** element

counts[100] does **not** exist

Why zero?

Because the index is an **offset** from the base address

Element 0 is 0 elements away from the start

An array in memory

address of element i = $\text{base} + i * \text{sizeof}(\text{element})$

For `uint32_t`, every element is 4 bytes apart

counts[0]	counts[1]	counts[2]	counts[3]	counts[4]	counts[5]
0x1F899BC	0x1F899C0	0x1F899C4	0x1F899C8	0x1F899CC	0x1F899D0

↑ **&counts[0] = base**

&counts[3] = base + 3 * 4 = 0x1F899C8

How many elements?

C:

`sizeof(counts) / sizeof(counts[0])`

C++ (std::array or std::vector):
`counts.size()`

Rust:

`counts.len()`

Out of bounds

C and C++: **no checking**

counts[100] reads or writes memory that is not yours

Undefined behavior - may crash, may silently corrupt

Rust: **checked at run time**

the program panics with a clear message and stops

This single difference is a large part of why Rust exists

Walking an array

C and C++:

```
for (int i = 0 ; i < 100 ; i++) {  
    psi[i] = adc_to_psi(counts[i]) ;  
}
```

Rust:

```
for i in 0..100 {  
    psi[i] = adc_to_psi(counts[i]) ;  
}
```

Arrays of structs

Put the two ideas together:

C, C++: `struct sample history[100] ;`
`history[7].psi = 50.0f ;`

Rust: `let mut history : [Sample ; 100] ;`
`history[7].psi = 50.0 ;`

Elements are `sizeof(struct sample)` bytes apart

Two dimensions

`uint8_t image[3][4] ;` 3 rows of 4

`let image : [[u8 ; 4] ; 3] ; Rust`

Stored one row after another (**row major**)

`image[1][2]` is at `base + (1*4 + 2)` bytes

It is still one flat block of memory

When the size is not known yet

C: malloc / free (later lecture)

C++: #include <vector>
 std::vector<uint32_t> counts ;
 counts.push_back(1456) ;

Rust: let mut counts : Vec<u32> = Vec::new() ;
 counts.push(1456) ;

Now Let Us Look Inside

You can now describe **structures** of memory

How do you **see** them while the program runs?

printf tells you what you decided to ask about
before you knew what was wrong

A **debugger** lets you stop and look at anything

What a debugger does

Loads your program but keeps **control** of it

Stops it where you choose (**breakpoints**)

Reads the memory of the stopped program

Uses the **symbol table** to turn addresses back into names

Shows a variable as its **type**, not as raw bytes

Debug information

The symbol table and type descriptions are placed in the executable by the compiler. The format is called **DWARF**

You have to ask for it:

```
gcc -g -O0 main.c      -o main
g++ -g -O0 main.cpp    -o main
cargo build             (dev profile has -g)
```

-O0 turns off optimization so the code matches your source

The tools

`gdb` - GNU debugger, Linux and embedded targets

`lldb` - LLVM debugger, the default on macOS

`rust-gdb` / `rust-lldb` - add Rust pretty printers

VS Code - a graphical front end for all of them

All of them do the same job with different spelling

Starting a session

```
gcc -g -O0 main.c -o main
```

```
gdb ./main
```

```
(gdb) break main
```

```
(gdb) run
```

The program stops at the first line of main
and hands control back to you

Everyday gdb Commands

Controlling the program:

break <file:line>	set a breakpoint
run	start the program
next	run one line, step over calls
step	run one line, step into calls
continue	run to the next breakpoint
finish	run until this function returns

Everyday gdb Commands

Looking around:

<code>print <expr></code>	show a value
<code>info locals</code>	show every local variable
<code>backtrace</code>	show the call stack
<code>list</code>	show the source
<code>quit</code>	leave the debugger

Short forms: `b r n s c bt p l q`

Looking at a variable

```
(gdb) print my_integer  
$1 = 55
```

```
(gdb) print/x my_integer  
$2 = 0x37
```

```
(gdb) print &my_integer  
$3 = (int64_t *) 0x7ffd1f899bc
```

Contents and address - the two ideas from lecture one

Looking at a struct

```
(gdb) print my_sample
```

```
$1 = {adc_counts = 1456, psi = 50, timestamp_ms = 12345}
```

```
(gdb) print my_sample.psi
```

```
$2 = 50
```

One command prints the **whole structure**, with names

What type is it?

(gdb) **p**type my_sample

```
type = struct sample {  
    uint32_t adc_counts;  
    float psi;  
    uint64_t timestamp_ms;  
}
```

The debugger knows this from the **DWARF** data

Looking at an array

(gdb) **print** counts

\$1 = {1456, 1442, 1460, 0 <repeats 97 times>}

(gdb) **print** counts[7]

\$2 = 0

One command prints **every element**

Part of an array: the @ operator

```
(gdb) print counts[2]@4  
$3 = {1460, 0, 0, 0}
```

Four elements, starting at index 2

Most useful when you have a pointer and a length:

```
(gdb) print *ptr@10
```

Looking at raw memory

(gdb) **x/16xb** &my_sample

x	examine
16	how many
x	format: hex (d, u, c, s, i)
b	size: byte (h, w, g)

This is where you can **see the padding bytes**

Who changed my variable?

```
(gdb) watch psi[3]
```

```
(gdb) continue
```

```
Hardware watchpoint 2: psi[3]
```

```
Old value = 0
```

```
New value = 50
```

```
update_sample () at main.c:42
```

The hardest bug to find by reading code

Debugging Rust

```
cargo build  
rust-gdb ./target/debug/myprog  
rust-lldb ./target/debug/myprog
```

The wrappers load **pretty printers**, so a Vec or a String prints readably instead of as internal pointers

Rust symbol names include the crate:
(gdb) break myprog::main

The same ideas in lldb

```
lldb ./main
```

```
(lldb) breakpoint set --name main    br s -n main
```

```
(lldb) run
```

```
(lldb) next / step / continue
```

```
(lldb) frame variable                fr v
```

```
(lldb) print my_sample
```

```
(lldb) memory read -f x -s 1 -c 16 &my_sample
```

Different spelling, identical ideas

The debugger in VS Code

Install the **C/C++** or **CodeLLDB** extension

For Rust: **rust-analyzer** plus CodeLLDB

Click in the gutter to set a breakpoint

F5 start, **F10** step over, **F11** step into

The VARIABLES pane

Lists every local variable, updated as you step

Expand a **struct** to see its members

Expand an **array** to see its elements

Hover over a name in the source to see its value

It is gdb or lldb underneath - the ideas still apply

A method for using it

1. Reproduce the failure
2. Set a breakpoint **before** the suspect line
3. Look at the variables - are they what you expect?
4. Step one line at a time and watch them change
5. The first value that **surprises** you is near the bug

What we can now describe

Variable	one name, one value, one address
Struct	different types, one block, fixed offsets
Array	one type repeated, $\text{base} + \text{index} * \text{size}$
Debugger	stop the program and look at all of it

Every one of these is still just **memory**

Let's use the debugger

Exercise: Using **structpractice**