

Scope and Persistence

Where We Left Off

Functions have parameters and local variables

Loops declare loop variables like `i`

Some names vanish at the closing brace
Some values survive between calls

Today: the rules behind both

Two Questions for Every Variable

Scope: where can the **name** be used?
a property of the source code

Persistence: how long does the **memory** exist?
a property of the running program

Also called **visibility** and **lifetime**

Names Are Addresses

A **variable** is a symbolic name for an address in the **data** section of the program

A **function** is a symbolic name for an address in the **instruction** section (the text)

You must be inside the **scope** to see the address as a symbol

Lexical Scoping

Scope defines where a variable or a function name is **valid**

A variable cannot be referenced by name from **outside** its scope

"Lexical" means it is decided by where the text appears in the source file

Scopes Are Braces

Lexical scopes are defined with `{ }`

```
{           outer scope
  {         inner scope
    }       inner names gone
  }         outer names gone
```

The same rule in C, C++, and Rust

Global Scope

The scope **outside of all functions**

C, C++: `int stepG ;`

Rust: `static STEP_G : i32 = 0 ;`

A variable defined there is a **global** variable

The **G** suffix marks it as global

Function Scope

The scope of a function's **parameters** is the complete body of the function

Local variables are in scope from their declaration to the closing brace

At the close of the scope, the symbol **disappears**

Scope Example in C (1 of 2)

```
void p ( char x )  
{  
    /* p, x */  
    char y ;  
    /* p, x, y */  
    char zoo ;  
    /* p, x, y, zoo */  
}  
/* p */  
double expG ;  
/* p, expG */
```

Scope Example in C (2 of 2)

```
void qsort ( char *s1, char *s2 )
{
    /* p, expG, qsort, s1, s2 */
    char b ;
    {
        int count ;
        /* p, expG, qsort, s1, s2, b, count */
    }
    /* p, expG, qsort, s1, s2, b */
}
/* p, expG, qsort */
```

The Same Idea in Rust

```
fn p ( x : u8 ) {  
    // p, x  
    let y : u8 = 0 ;  
    // p, x, y  
    {  
        let count = 0 ;           // p, x, y, count  
    }  
    // p, x, y  
}
```

Order Matters in C, Not in Rust

C, C++: a name is visible from its declaration **downward** to the end of its scope

Rust: functions, statics, and types are visible in the **whole module**, even above them

Rust local `let` variables still follow top-to-bottom order

Shadowing

An inner scope may reuse an outer name:

```
int x = 5 ;  
{ int x = 10 ; }
```

inner x hides outer x

Rust also allows it in the **same** scope:

```
let x = 5 ;  
let x = x * 2 ;
```

a new x, now 10

Scope and Loops

```
for ( int i = 0 ; i < n ; i++ ) { ... }  
    i is gone here
```

```
for i in 0..n { ... }  
    i is gone here too
```

A loop variable's scope is the loop

But Where Is the Memory?


A variable definition allocates memory

Where it goes depends on the variable's **scope** and its **lifetime**

Three places memory can live:

data segment stack heap

Memory Map of a Linux Process



Region 1: The Data Segment

Persistent variables live here

Size known at **compile time** - part of the object code

Data: variables with an initial value

BSS: variables with no initial value
set to **zero** when the program loads

Region 2: The Stack


Holds each call's **arguments** and **local** variables

A **stack** data structure: push on call, pop on return

Located by the **stack pointer** and **frame pointer**

Memory on the stack is **uninitialized** and **temporary**

A Stack Frame



Region 3: The Heap

Dynamic memory - size decided at run time

The **program** decides when it is created
and when it is released

Lives until it is freed - not tied to a
function call

The Heap in C

```
float *buf = malloc ( 100 * sizeof ( float ) ) ;  
if ( buf == NULL ) { /* out of memory */ }
```

... use buf[0] through buf[99] ...

```
free ( buf ) ;
```

Also: calloc (zeroed), realloc (resize)

The Heap in C++

```
float *buf = new float[100] ;
```

```
...
```

```
delete[] buf ;
```

Better: let a type own the memory

```
std::vector<float> buf ( 100 ) ;
```

```
auto p = std::make_unique<sample_t> ( ) ;
```

Freed automatically at the closing brace

The Heap in Rust

```
let buf = vec![ 0.0f32 ; 100 ] ;  
let s = Box::new ( sample ) ;
```

No free. No delete.

When the **owner** goes out of scope,
the heap memory is released

Heap Bugs in C

Leak: forget to free - memory is never returned

Use after free: read memory already given back

Double free: free the same memory twice

C compiles all three without complaint

Rust rejects all three at compile time

Automatic Variables

Declared inside a function

Created on the **stack** when the call starts
Gone when the call returns

Called **auto** variables in C: automatic allocation and release

Auto Variables Start as Garbage

```
void my_function ( int expand )  
{  
    int step ;  
    step = step + expand ; bug  
}
```

step was never initialized

C, C++: the program **does not crash** - it
silently uses whatever was on the stack

Rust Refuses to Compile It

```
fn my_function ( expand : i32 ) {  
    let step : i32 ;  
    let total = step + expand ;  
}
```

error: used binding `step` isn't initialized

Rust checks every path before any use

Persistent Variables

Their memory exists for the **entire** program execution

Stored in the **data segment**

Initialized to **zero** or to a value you give, **once**, before main runs

Persistent Variable Visibility

The visibility of a persistent variable is determined by its **scope**:

inside a function	local static
outside, file only	file static
outside, all files	global

Same lifetime - three different scopes

Local Static

```
void my_function ( int expand )  
{  
    static int stepL ;  
    stepL = stepL + expand ;  
}
```

Visible only **inside** my_function
But the value **persists** between calls

Example: Counting Calls

```
int count_calls ( void ) {  
    static int countL = 0 ;  
    countL = countL + 1 ;  
    return countL ;  
}
```

Returns 1, then 2, then 3, ...

The `= 0` happens **once**, not on every call

Local Static in Rust

```
use std::sync::atomic::{ AtomicU32, Ordering } ;

fn count_calls ( ) -> u32 {
    static COUNT_L : AtomicU32 = AtomicU32::new ( 0
) ;
    COUNT_L.fetch_add ( 1, Ordering::Relaxed ) + 1
}
```

A mutable static must be made **safe to share**

Declarations

Used to reference a variable or a function
before it is defined

Say what it **is** - its type

Apply to functions (instruction text)
and to variables (memory addresses)

Definitions

Describe the **implementation** of a function or variable

Function definition: the parameters and the statements that implement it

Variable definition: describes **and allocates** space for the address

Function Declaration

Known as a **prototype**

```
extern <return_type> <name> ( arglist ) ;
```

```
extern int access_data ( char *array,  
    int idx, int num_elements ) ;
```

Takes a pointer, an index, and a count
Returns an int

extern on Functions

extern extends the visibility of a function to other files

Not strictly necessary for functions,
but recommended

Many **identical** declarations are allowed
Only **one** definition

Function Definition

```
int access_data ( char *array, int idx,  
                 int num_elements ) {  
    if ( ( idx >= 0 ) && ( idx < num_elements ) ) {  
        return array[idx] ;  
    }  
    return -1 ;           // out of bounds  
}
```

This becomes **text**: machine instructions

extern on Variables

```
extern double myNumG ;
```

Describes the **type** but does **not** allocate

A placeholder for the real address

The address is **resolved** at link time

myNumG is 8 bytes of double - defined elsewhere

Variable Definition

```
double myNumG ;
```

myNumG is the starting address of 8 bytes holding a double

This statement **allocates** the memory

Exactly **one** file may contain this line

Declarations in C++

The same rules as C: prototypes, extern

Declarations go in a **header** (.h / .hpp)

Definitions go in a **source** file (.cpp)

Break the **One Definition Rule** and the linker reports a duplicate symbol

Declarations in Rust

Rust has **no** separate declarations
and **no** header files

```
mod sensor ;           load sensor.rs as a module
use sensor::adc_to_psi ;   bring the name into
scope
```

The compiler sees the whole crate at once

Global Variables

```
int stepG ;
```

```
int main ( int argc, char *argv[] )  
{  
    ...  
}
```

Four bytes in the **data** section,
for the **lifetime** of the program

Using a Global from Another File

another.c:

```
extern int stepG ;
```

```
void reset ( void ) {  
    stepG = 10 ;  
}
```

Usable in any scope and file that declares it

Globals in C++

Same as C, plus **namespaces**:

```
namespace sensor {  
    int stepG ;  
}
```

```
sensor::stepG = 10 ;
```

A namespace keeps global names from colliding

Globals in Rust

```
static mut STEP_G : i32 = 0 ;
```

```
unsafe { STEP_G = 10 ; }
```

Every read or write of a **mutable** global
needs **unsafe**

Rust is telling you globals are risky

Safe Globals in Rust

```
use std::sync::atomic::{ AtomicI32, Ordering } ;  
  
static STEP_G : AtomicI32 = AtomicI32::new ( 0 ) ;  
  
STEP_G.store ( 10, Ordering::Relaxed ) ;  
let s = STEP_G.load ( Ordering::Relaxed ) ;
```

No unsafe - the type makes sharing safe

Limiting Persistent Visibility

What if a persistent variable should be visible to **one file** only?

Use a **static** definition:

```
static int stepS ;
```

What is its scope?

File Scope: static

Defined outside any function

Scope: every function that **follows** in the file

```
static int stepS ;
```

```
void my_function ( int expand ) {  
    stepS = expand ;  
}
```

Other files cannot reach it, even with extern

File Scope in C++

static works as in C

Preferred: an **unnamed namespace**

```
namespace {  
    int stepS ;  
}
```

Everything inside is private to this file

Module Scope in Rust

Everything is **private** to its module
unless marked **pub**

```
static STEP_S : i32 = 0 ;    this module only  
pub static STEP_G : i32 = 0 ;    visible to other  
modules
```

Rust's default is C's `static`

Limiting Function Visibility

A function can be private to its file too

Declaration:

```
static <return_type> <name> ( arglist ) ;
```

Definition:

```
static <return_type> <name> ( arglist ) {  
    ...  
}
```

A Private Function

file: compute.c

```
static int my_static_func ( int x ) {  
    return ( x * x ) - 1 ;  
}
```

This function only exists in **compute.o**
No other file can call it

Private Functions in C++ and Rust

C++:

static or an unnamed namespace

private: member functions of a class

Rust:

fn helper () private to the module

pub fn api () visible to other modules

pub(crate) fn tool () visible in this crate

only

Global and Static: main.c

```
#include <stdio.h>
extern void my_sub ( void ) ;
int countG ;                                global definition
static int resetS ;                         static definition

int main ( int argc, char *argv[] ) {
    countG = 1 ;
    resetS = 1 ;
    my_sub ( ) ;
    fprintf ( stderr, "countG is %d not 1\n", countG ) ;
    fprintf ( stderr, "reset is still %d\n", resetS ) ;
    return 0 ;
}
```

Global and Static: mysub.c

```
#include <stdio.h>
extern int countG ;
static int resetS ;
```

global reference
a **different** resetS

```
void my_sub ( void ) {
    countG = 2 ;
    resetS = 0 ;
    fprintf ( stderr, "countG in my_sub is %d\n", countG ) ;
    fprintf ( stderr, "resetS in my_sub is %d\n", resetS ) ;
}
```

The Symbol Tables

```
gcc -c main.c ; nm main.o  
mysub.o
```

```
C _countG  
T _main  
U _my_sub  
s _resetS
```

```
gcc -c mysub.c ; nm
```

```
U _countG  
T _my_sub  
s _resetS
```

T text **U** undefined - linker must find it
C global data **s** lowercase = file private

After Linking

```
gcc -o myprog main.o mysub.o ; nm myprog
```

S	_countG	one common copy
s	_resetS	
s	_resetS	two separate copies
T	_main	
T	_my_sub	

The linker joins globals by name; statics stay apart

Running It: Side Effects

./myprog

countG in my_sub is 2

resetS in my_sub is 0

countG is 2 not 1

side effect

reset is still 1

no side effect

main set countG to 1; my_sub changed it to 2

Why Globals Are Slow

```
sw  $2, %gp_rel(countG)($28)  countG = 1  
jal my_sub                    call  
lw  $2, %gp_rel(countG)($28)  reload countG
```

The compiler must assume **any** call it cannot see
may change a global - so it reloads it

Prefer local variables and const

The Same Program in Rust

main.rs

```
mod mysub ;  
pub static COUNT_G : AtomicI32 = AtomicI32::new  
( 0 ) ;  
static RESET_S : AtomicI32 = AtomicI32::new ( 0 ) ;
```

mysub.rs

```
use crate::COUNT_G ;  
static RESET_S : AtomicI32 = AtomicI32::new ( 0 ) ;
```

Two RESET_S - one per module, like C's static

The volatile Attribute

Tells the compiler the variable could
change at any time

The compiler will not optimize its reads
or writes away

Declared in C and C++ with **volatile**

When You Need volatile

In practice, only three kinds of variable:

1. Memory-mapped **peripheral registers**
2. Globals changed by an **interrupt** routine
3. Globals shared by **multiple tasks**
in a multi-threaded program

Register Without volatile

```
uint8_t *my_reg = ( uint8_t * ) 0xC000 ;
```

```
// wait for the register to read non-zero  
do {  
} while ( ! ( *my_reg ) ) ;
```

The optimizer reads `*my_reg` **once**,
before the loop. It sees nothing that could
change it - so the loop never ends

Register With volatile

```
uint8_t volatile *my_reg = ( uint8_t * ) 0xC000 ;
```

```
do {  
} while ( ! ( *my_reg ) ) ;
```

Now every pass **re-reads** the register

The **data** is volatile, not the pointer

Interrupt Flag Without volatile

```
bool gb_etx_foundG = false ;
```

```
main:      while ( !gb_etx_foundG ) { }
```

```
interrupt: if ( ETX == rx_char )  
           gb_etx_foundG = true ;
```

No call in the loop - the optimizer may load the flag once and loop forever

Interrupt Flag With volatile

```
bool volatile gb_etx_foundG = false ;
```

```
main:      while ( !gb_etx_foundG ) { }
```

```
interrupt: if ( ETX == rx_char )  
           gb_etx_foundG = true ;
```

volatile guarantees the compiler reads the flag on every pass

volatile in Rust

Rust has **no** volatile keyword

Registers: explicit volatile access

```
let my_reg = 0xC000 as *const u8 ;  
while unsafe { my_reg.read_volatile ( ) } == 0  
{ }
```

Embedded register crates wrap this for you

Interrupt Flags in Rust

```
static ETX_FOUND_G : AtomicBool =  
    AtomicBool::new ( false ) ;
```

```
main:      while !ETX_FOUND_G.load  
( Ordering::Acquire ) { }
```

```
interrupt: ETX_FOUND_G.store ( true,  
    Ordering::Release ) ;
```

An atomic is both safe and never optimized away

volatile Is Not Enough for Threads

volatile stops the compiler removing accesses

It does **not** make $x = x + 1$ indivisible

Two threads can still lose an update

Use atomics for shared data:

C: `<stdatomic.h>` C++: `std::atomic`

Rust: `AtomicI32`, `AtomicBool`, ...

Scope and Lifetime Together

Kind	Scope	Lifetime	Lives in
auto (local)	the block	the call	stack
local static	the function	whole program	data
file static	the file	whole program	data
global	every file (extern)	whole program	data
heap	via a pointer	until freed	heap

Across the Three Languages

	C	C++	Rust
global	<code>int xG ;</code>	<code>int xG ;</code> or <code>namespace</code>	<code>pub static</code> (atomic)
file private	<code>static int xS ;</code>	<code>namespace { }</code>	default (no pub)
local static	<code>static int xL ;</code>	<code>static int xL ;</code>	<code>static</code> (atomic)
declare	<code>extern</code>	<code>extern</code>	<code>mod + use</code>
heap	<code>malloc / free</code>	<code>vector, unique_ptr</code>	<code>Vec, Box</code>
volatile	<code>volatile</code>	<code>volatile</code>	<code>read_volatile</code>

Good Habits

Prefer **local** variables and parameters

Give every name the **smallest** scope that works

Make file-only data **static** (C) or private (Rust)

Mark names: **G** global, **S** file static, **L** local static

Use **const** for anything that should not change

What We Can Now Do

Scope	where a name is visible
Lifetime	how long its memory exists
Regions	data segment, stack, heap
Keywords	extern static volatile pub

Let's practice

Exercise: Using **scopepractice**